

# ELEN4020A - Data Intensive Computing

Lecturer: Prof E. J. Otoo

## Laboratory 3: Using MapReduce Framework

16 April 2019

Tyson Cross 1239448 Michael Nortje 1389486 Josh Isserow 675720

### Abstract

MapReduce techniques were used to process several text files of different lengths with word-count and indexing algorithms in Python with the mrjob 0.6.8 framework for the Hadoop 2.9.2 libraries. Three algorithms were developed: to return the number of occurrences of each word in a text; to return the top-K frequencies of word occurrence in a text; and to return an inverted line index for each word in a text.

## 1 INTRODUCTION

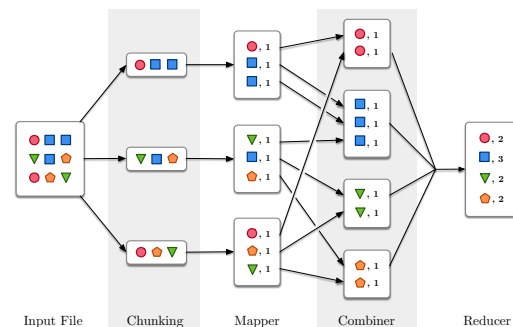
A recent projection of the global growth in the Big Data market has forecast the segment to exceed a valuation of \$70 billion by 2022, almost doubling the current market size [1]. The use of very large data sets is central to the future plans of many major corporations and international organisations, with the quantity of global big data predicted to rise to above 40,000 Exabytes by 2020 [2].

One of the most important tools to help process these very large data sets is MapReduce, a framework capable of processing copious amounts of data in a fault-tolerant method [3] using implementations such as the Apache Hadoop framework, written in Java. MapReduce can utilise massive parallel clusters for data processing, accessing terabyte and petabyte level of data [3]. The principal behind MapReduce is to utilise an uncomplicated library capable of parallelization, distribution of data and load-balancing.

This report explains the use of a framework to implement of MapReduce for three algorithms: to return the number of occurrences of each word in a text; to return the top-K frequencies of word occurrence in a text; and to return an inverted line index for each word in a text.

## 2 MAPREDUCE

MapReduce is central to the Apache Hadoop library [3]. A diagram showing a high level overview of the data-flow steps in MapReduce is shown in Figure 1. There are two basic methods of operation: `map` and `reduce`. `Map` takes a single key and value pair as input, to produce a (key,values) output (the mapper can also return nothing, if no values for the input key are found). `Reduce` uses the output produced by `Map` and amalgamates the key/values pairs.



**Fig. 1: MapReduce diagram**

There is also an optional extra optimisation step, `combiner` which takes a key, and subset of values, and returns none or more key/values pairs, to decrease the overall amount of data passing through the system before being handed to the reducer. MapReduce also can include partitioners, that split up the key space before reduction, and counters which can be used to provide statistics

for analysis of performance of the mapper and reducer operations.

When a program conforms to the MapReduce standard (such as Hadoop), parallelization and execution are managed from the controlling node, and carried out by the workers, which can be affordable commodity machines rather than specialised cluster computers [4], [5]. Run-time systems assume responsibility of separating input data, arranging execution of the programs, managing failures experienced by the machines as well inter-machine communication. This simplifies the ability to utilise large distributed systems without prior parallel system experience [4]. Part of the Hadoop framework is the Hadoop Distributed File System (HDFS), which is intended to be highly fault-tolerance, assuming that hardware failure involving large networks of many machines can occur at any time. The framework used for this laboratory, mrjob 0.6.8 supports multiple non-HDFS systems.

The MapReduce framework consists of two basic components: a main controlling node (either a "job tracker" or "resource manager"), and the number of worker nodes, the controller schedules and directs the jobs, which are a series of tasks (or steps) executed by the workers [6].

### 2.1 mrjob

Mrjob is an Apache Licensed (Version 2.0) Hadoop framework which allows for easy routing for programs in Python that run on Hadoop [6]. The framework allows for easy writing and management of MapReduce jobs. With mrjob, the jobs are written as classes that execute functions according to a `step` order, controlling of the flow of execution.

Mrjob is well documented, and lets the user run code in Python instead of Java without the need for directly controlling or installing Hadoop [6]. This allows for writing and running tests for MapReduce using the Hadoop streaming back-end. The framework offers a consistent interface for all environments that are supported - i.e. the Python code itself does not change depending on the operating system or hardware that it is run on. Mrjob offers easy debugging through the use

of local error trace-backs in the console, rather than in a log file. Part of mrjob's ease of use is that it automatically handles serialization and de-serialization of input and output key/value pairs. This removes the need for manually loading and dumping `json` objects.

One disadvantage of using mrjob is actually its ease of use. Since it is greatly simplified, it does not have the same access to the Hadoop API that a lower-level framework might offer.

## 3 ALGORITHMS

MapReduce was designed and intended to be highly scalable, able to perform data operations on a local computer, or across very large clusters of multiple machines with their own storage. For the laboratory code execution a single 4-core laptop was used. All algorithms were implemented in Python 3.7 using the mrjob API and the Hadoop back-end. The inputs for the laboratory are plain-text files, of varying lengths and line counts. An additional input file required is a stop-words text file, which lists all the words that should be discounted. Stop words are commonly used to filter out common joining words so that only the data considered useful remains [7].

### 3.1 Word Count

A simple word count algorithm providing the number of occurrences of all words in a given text was implemented. The implemented code is listed in Code B1.

The algorithm is implemented as a class `WordCount`, which inherits from the principle `MRJob` class. The steps of the algorithm are as follows: first an initialisation mapper step (`mapper_init`) loads and separates out the stop words onto an property on the main class. The input text file is automatically split up into chunks, and distributed to the mappers before the main mapper step, shown as pseudocode in Algorithm 1:

```

Input: line
Output: word, 1
for each word in line (stripped of punctuation)
do
    make word lower case
    if word not in stop_words then
        return word, 1
    end
end

```

#### Algorithm 1: Word Count: mapper

Each mapper finds a single unique instance of a word in a line (in the backend, storing an internal local and total byte-offset to determine if a word has been considered yet). The next step creates combiners, which sort the summed sub-values (word occurrences per mapper per line) received for each key, which are then passed to the reducer step, which sums the total word occurrences across the entire input text. These value are sorted by mrjob and Hadoop streaming, before printing the final steps' generator values to `std.out` as a list of words and their count in the text.

A command-line argument is included for the python script, to allow for providing a required external custom stop-words file. The file can be empty. By default, the stop words file listed in Table A2 was used for all of the processing of the texts, except for *Hamlet*, which used the list of words shown in Table A3, due to the Elizabethan language in the play.

The mapper functions are responsible for processing the words from the input text file one line at a time. The combiner function is responsible for combining sub-values associated with each key. The reducer functions perform the addition of the values for each key of numbers, yielding the final summed (key,value) pairs. In order to control the individual task steps, the mrjob scheduler function `MRStep` was overridden, with manually defined mapper, combiner and reducer functions.

### 3.2 Top K Word Frequencies

This algorithm returns a list of the K most frequently occurring words in a given text, excluding stop words, for  $K = 10$  and  $K = 20$ . The implemented code is shown in Code B2. The steps of the algorithm are similar to the previous word count, only with an extra reducer step. In this final

reducer stage, the last and current frequency of word incidences are kept track of inside a nested conditional statement within a *for* loop, which traverses through the sorted word/counts pairs, and stops returning values once K values have been returned. This step is just a post-process step to restrict the computed key/values being returned, and altering K has no affect on the asymptotic behaviour of the algorithm. Because the entire text has to be counted (and so every word must be considered) it is unclear how K could be used to more substantially shorten the number of execution steps required, given the parallel process of computation.

```

Input: word_counts
K ← user input
i ← 0
last_freq ← 0
current_freq ← 0
for each count and key in word_counts (in
    descending order) do
    if i equal to K then
        | current_freq ← count
    else if i less than K then
        last_freq ← current_freq
        current_freq ← count
        if last_freq not equal to
            | current_freq then
                | i ← i+1
        end
    end
end

```

#### Algorithm 2: Word Frequency: reducer\_sort\_counts

The frequency limit was implemented as a command-line argument, which defaults to the top 10 if not provided explicitly. The algorithm returns the number of most common frequencies, which can be more or less than the number of words returned. This is because more than one word can occur at the same frequency of occurrence, or there could be fewer words than the requested number of frequencies. Because the entire text must be processed and then sorted before determining the most common frequencies, there is no performance or speed difference irrespective of the number of top frequencies requested.

The top 10 word frequency results for the test texts are shown in Figure A.1f.

### 3.3 Inverted Line Index

This algorithm provides the line numbers where a particular word can be found. Each individual mapper only receives an individual line of text, with no idea of where in the text it was located. When returning the processed key (word) and value (line), it quickly becomes difficult for a human to read the results, particularly with long line lengths. The implemented code is shown in Code B3.

One solution to contextualise the line index, while still returning each entire line of text with each word, is to pre-process the input text, pre-pending each line with a number. The provided text for the laboratory, `File2ForLab3.txt` (a Gutenberg plain-text copy of Plato's *Republic*) has already been pre-processed in this manner. For the testing in this report, multiple additional texts were used. For brevity and readability, instead of returning the entire line of text pre-pended with a line-number, it was decided to return only the line number, and the integrate the pre-processing programatically.

```
Input: word_counts
for each num and line in f do
  if i greater than or equal to K then
    break
  for each word in line do
    word ← lowercase word
    if word not a stopword then
      yield word and num
    end
  end
end
```

**Algorithm 3:** Word Inverse Index: mapper\_raw

To accomplish this, the `mapper_raw` function was used, which provides a path to the entire file, available for each mapper. The pseudocode for this function is shown in Algorithm 3. Mrjob and Hadoop take care of all the low level data processing, providing the correct locations, lines and accessible local copies of the input file. The mapper first enumerates each line, and then

performs the word processing. The function then yields a generator for each line, returning the key (word) and value (line number). This is much easier for a human to read the final returned result. A optional command-line argument was added to limit how many lines of the file to process, and another option to limit the number of indexed lines that are listed per word in the output.

## 4 PERFORMANCE EVALUATION

Several differently sized tests from the Gutenberg project were used as data for processes, listed in Table A1. The files were input as UTF-8 encoded plain-text files, with the Gutenberg metadata stripped. The timing of the results is shown in Figure A.2. Timing was performed using the `datetime` function, and output using a call to `sys.stderr.write` to output the delta time to the console.

The Inverse Index algorithm is extremely rapid when asked to only process 50 lines, returning an unsorted result in less than 165 ms, with an average of 150 ms (for all texts). When the entire text is processed, with no limit of the number of indexed lined returned per word, the shortest text *Symbols and Signs* with 200 lines returned a result in 208 ms. The longest test, *In Search of Lost Time* (comprised of all 7 volumes of the English translation of Marcel Proust's *À la recherche du temps perdu* at 116365 lines) returns the indexed results in 8421 ms. The most common word, 'time', occurs on 2695 lines in the text.


## 5 CONCLUSION

A variety of common data tasks was processed using the MapReduce framework. A word count algorithm, K-query algorithm and inversion index of text algorithm were written and evaluated, using Python and the mrjob framework. The algorithms were run on a variety of short, medium and long texts. The performance of the algorithms were timed, and the results were compared and discussed.


## REFERENCES

- [1] Wikibon.com, “Big data market revenue forecast worldwide 2011-2027,” <https://www.statista.com/statistics/254266/global-big-data-market-forecast/>, March 2019, (Accessed on 04/15/2019).
- [2] D. Srivastava and D. Sinwar, “Statistical analysis of big data using hadoop: A review,” vol. 5, 01 2018.
- [3] S. N. Khezr and N. J. Navimipour, “Mapreduce and its applications, challenges, and architecture: a comprehensive review and directions for future research,” *Journal of Grid Computing*, vol. 15, no. 3, pp. 295–321, 2017.
- [4] J. Dean and S. Ghemawat, “Mapreduce: simplified data processing on large clusters,” *Communications of the ACM*, vol. 51, no. 1, pp. 107–113, 2008.
- [5] S. Maitrey and C. Jha, “Mapreduce: simplified data analysis of big data,” *Procedia Computer Science*, vol. 57, pp. 563–571, 2015.
- [6] Yelp and Contributors, “MRJob v0.6.8.dev0 API Documentation,” <https://mrjob.readthedocs.io/>, January 2019, (Accessed on 04/15/2019).
- [7] A. Rajaraman and J. D. Ullman, *Mining of massive datasets*. Cambridge University Press, 2011.


## APPENDIX A




(a) Top 10 Word Frequencies in 'Hamlet'




(b) Top 10 Word Frequencies in 'The Third Policeman'




(c) Top 10 Word Frequencies in 'The Iliad'



(d) Top 10 Word Frequencies in 'The Republic'




(e) Top 10 Word Frequencies in 'In Search of Lost Time'



(f) Top 1000 Words Frequencies in 'In Search of Lost Time'

**Fig. A.1:** Main Caption



**Fig. A.2:** Algorithm Timings

**TABLE A1:** Texts Used in Testing

Text	Author	Words	Lines	Characters
<i>Symbols and Signs</i>	Vladimir Nabokov	2230	206	12580
<i>Hamlet</i>	William Shakespeare	32065	4460	191726
<i>The Third Policeman</i>	Flann O'Brien	74728	3729	403872
<i>Iliad</i>	Homer	153320	13440	808161
<i>The Republic</i>	Plato	217340	24295	1195060
<i>In Search of Lost Time</i>	Marcel Proust	1361977	116365	7498967

**TABLE A2: stop\_words.txt**

[numerals 0-100] [numerals for dates 1990-2020] [punctuation marks] ' 'i 'll a about above across after afterwards again against all almost alone along already also although always am among amongst amount an and another any anyhow anyone anything anyway anywhere are around as at b back be because been before beforehand being below beside besides between both bottom but by c call can cannot cant co come con could couldnt d de did didn't do does doesn't doing don don't done down due during e each eg eight either eleven else elsewhere enough etc even ever every everyone everything everywhere except f few fifteen fifty fill find fire first five for former formerly forty found four from front full further g get give go h had has hasnt have having he hence her here hereafter hereby herein hereupon hers herself him himself his how however hundred i i'll ie if in inc indeed into is it its itself j just k keep l last latter latterly least less like ll ltd m made many may me meanwhile might mill mine mme more moreover most mostly move much must my myself n name namely neither never nevertheless next nine no nobody none noone nor not nothing now nowhere o of off often on once one only onto or other others otherwise our ours ourselves out over own p part per perhaps please put q r rather re s said same say see seem seemed seeming seems serious several she should since six sixty so some somehow someone something sometime sometimes somewhere still such system t take ten than that the their theirs them themselves then there thereafter thereby therefore therein thereupon these they thing third this those though three through throughout thru thus to together too toward towards twelve twenty two u un under until up upon us v very via w was we we'll well were what whatever when whenever where whereafter whereas whereby wherein whereupon wherever whether which while who whoever whole whom whose why will with within without would x y yet you your yours yourself yourselves z

**TABLE A3: stop\_words\_shakespeare.txt**

[numerals 0-100] [numerals for dates 1990-2020] [punctuation marks] 'tis 'a 'a' 'and 'as 'but 'by 'for 'he 'if 'in 'it 'o 's 'seems 'should' 'so 't 'that 'there 'thine 'thus 'tis 'to 'twas 'tween 'twould 'we 'well 'while 'would a ' about above across after afterwards again against all almost alone along already also although always am among amongst amoungst amount an and another any anyhow anyone anything anyway anywhere are around art as at ay b back be because been before beforehand being below beside besides between both bottom but by c call can cannot cant claud co come con could couldnt d de did didn't do does doesn't doing don don't done dost doth down due during e each eg eight either eleven else elsewhere enough enter ere etc even ever every everyone everything everywhere except exeunt exit f few fifteen fifty fill find fire first five for former formerly forty found four from front full further g ger get give go h had ham has hasnt hast hath have having he hence hence her here hereafter hereby herein hereupon hers herself him himself his hither ho hor how however hundred i i' i'll ie if in inc indeed into is it its itself j just k keep l laer last latter latterly least less let like ll ltd m made many may me meanwhile might mill mine more moreover most mostly move much must my myself n name namely neither never nevertheless next nigh nine no nobody none noone nor not nothing now nowhere o o'er of off oft often on once one only onto oph or other others otherwise our ours ourselves out over own p part per perhaps please pol put q r rather re ros s same see seem seemed seeming seems serious several shall she should since six sixty so some somehow someone something sometime sometimes somewhere st still such system t take ten th' than that the thee their theirs them themselves then thence there thereafter thereby therefore therein thereupon these they thine thing third this thither those thou though three through throughout thru thus thy tither to together too toward towards twelve twenty two u un under until up upon us v very via w was wast we we'll well were what whatever when whence whenever where whereafter whereas whereby wherefore wherein whereto whereupon wherever whether which while whither who whoever whole whom whose why will with withal within without would x y ye yet yon yonder you your yours yourself yourselves z



## APPENDIX B

### Code B1: WordCount.py

```

1  from mrjob.job import MRJob
2  from mrjob.step import MRStep
3  from datetime import datetime
4  import sys
5  import re
6
7  WORD_RE = re.compile(r"[\w']+")
8
9  class WordCount(MRJob):
10     '''Returns the number of occurrences of words in the input text'''
11
12     FILES = ['stop_words.txt']
13     SORT_VALUES = True
14
15     def configure_args(self):
16         super(WordCount, self).configure_args()
17         self.add_file_arg(
18             '--stop-words',
19             metavar='STOP_WORDS_FILE',
20             dest='stop_words',
21             type=str,
22             default='stop_words.txt',
23             help='Input stop words text file ')
24
25     def steps(self):
26         return [
27             MRStep(
28                 mapper_init=self.mapper_init,
29                 mapper=self.mapper_get_words,
30                 combiner=self.combiner,
31                 reducer=self.reducer,
32             ),
33             MRStep(
34                 reducer=self.reducer_sort_counts,
35             )
36         ]
37
38     def mapper_init(self):
39         with open(self.options.stop_words) as f:
40             self.stop_words = set(line.strip() for line in f)
41
42     def mapper_get_words(self, _, line):
43         for word in WORD_RE.findall(line):
44             word = word.lower()
45             if word not in self.stop_words:
46                 yield (word.lower(), 1)
47
48     def combiner(self, word, counts):
49         yield (word, sum(counts))
50
51     def reducer(self, key, values):
52         yield None, (sum(values), key)
53
54     def reducer_sort_counts(self, _, word_counts):
55         for count, key in sorted(word_counts, reverse=True):
56             yield (count, key)
57
58 if __name__ == '__main__':
59     start_time = datetime.now()
60     WordCount.run()
61     end_time = datetime.now()
62     elapsed_time = (end_time - start_time)*1000
63     sys.stderr.write("Total Seconds WordCount.py: {0} microseconds\n".format(elapsed_time.total_seconds()))

```

```

1  from mrjob.job import MRJob
2  from mrjob.step import MRStep
3  from datetime import datetime
4  import sys
5  import re
6
7  WORD_RE = re.compile(r"[\w']+")
8
9  class WordFreq(MRJob):
10     ''' This function returns the top K frequency occurrences of words in a
11         provided text.
12         The number of words may be greater or less than K, depending if there
13         are several words with
14         the same number of occurrences in the text, or if there are fewer
15         words than K in the text. '''
16
17     FILES = ['stop_words.txt']
18     SORT_VALUES = True
19
20     def configure_args(self):
21         super(WordFreq, self).configure_args()
22         self.add_passthru_arg(
23             '--limit',
24             metavar='K',
25             dest='K',
26             type=int,
27             default=10,
28             help='Input stop words text file ')
29         self.add_file_arg(
30             '--stop-words',
31             metavar='STOP_WORDS_FILE',
32             dest='stop_words',
33             type=str,
34             default='stop_words.txt',
35             help='Number of highest occurrences to return')
36
37     def steps(self):
38         return [
39             MRStep(
40                 mapper_init=self.mapper_init,
41                 mapper=self.mapper_get_words,
42                 combiner=self.combiner,
43                 reducer=self.reducer,
44             ),
45             MRStep(
46                 reducer=self.reducer_sort_counts,
47             )
48         ]
49
50     def mapper_init(self):
51         with open(self.options.stop_words) as f:
52             self.stop_words = set(line.strip() for line in f)
53
54     def mapper_get_words(self, _, line):
55         for word in WORD_RE.findall(line):
56             word = word.lower()
57             if word not in self.stop_words:
58                 yield (word.lower(), 1)
59
60     def combiner(self, word, counts):
61         yield (word, sum(counts))
62
63     def reducer(self, key, values):
64         yield None, (sum(values), key)
65
66     def reducer_sort_counts(self, _, word_counts):
67         K = self.options.K

```

```

65     i = 0
66     last_freq = 0
67     current_freq = 0
68     for count, key in sorted(word_counts, reverse=True):
69         if (i==K):
70             current_freq = int(count)
71             if (i < K):
72                 last_freq = current_freq
73                 current_freq = int(count)
74                 if (last_freq != current_freq):
75                     i+=1
76             yield (count, key)
77
78 if __name__ == '__main__':
79     start_time = datetime.now()
80     WordFreq.run()
81     end_time = datetime.now()
82     elapsed_time = (end_time - start_time)*1000
83     sys.stderr.write("Total Seconds WordFreq.py: ({0}) microseconds\n".format(elapsed_time.total_seconds()))

```

#### Code B3: WordInverseIndex.py

```

1  from mrjob.job import MRJob
2  from mrjob.step import MRStep
3  from datetime import datetime
4  import itertools
5  import sys
6  import re
7
8  WORD_RE = re.compile(r"[\w']+")
9
10 class WordInverseIndex(MRJob):
11     ''' Returns the inverse index of the words in the first K lines in the
        input text'''
12
13     FILES = ['stop_words.txt']
14     SORT_VALUES = True
15
16     def configure_args(self):
17         super(WordInverseIndex, self).configure_args()
18         self.add_file_arg( '--stop-words',
19                           metavar='STOP_WORDS_FILE',
20                           dest='stop_words',
21                           type=str,
22                           default='stop_words.txt',
23                           help='Input stop words text file ')
24         self.add_passthru_arg( '--limit',
25                                metavar='K',
26                                dest='K',
27                                type=int,
28                                default=-1,
29                                help='Number of lines of the text to process')
30         self.add_passthru_arg( '--index-limit',
31                                metavar='L',
32                                dest='L',
33                                type=int,
34                                default=-1,
35                                help='Number of indexed lines per word to return')
36
37     def steps(self):
38         return [
39             MRStep(
40                 mapper_init=self.mapper_init,
41                 mapper_raw=self.mapper_raw,
42                 reducer=self.reducer,
43             ),
44             MRStep(
45                 reducer=self.reducer_sort_counts,
46             )
47         ]

```

```

48
49 def mapper_init(self):
50     with open(self.options.stop_words) as f:
51         self.stop_words = set(line.strip() for line in f)
52
53 def mapper_raw(self, path, _):
54     with open(path, 'r') as f:
55         for num, line in enumerate(f, 1):
56             if (self.options.K != -1) & (num >= self.options.K): break
57             for word in WORD_RE.findall(line):
58                 word = word.lower()
59                 if word not in self.stop_words:
60                     yield (word.lower(), num) # optionally, can pass [num, line.
61                                             rstrip()]
62
63 def reducer(self, key, values):
64     if (self.options.L == -1):
65         lines_list = list(values)
66     else:
67         lines_list = list(itertools.islice(values, self.options.L)) #
68     yield None, (key, ','.join(str(line) for line in lines_list))
69
70 def reducer_sort_counts(self, _, values):
71     for key, lines in sorted(values):
72         yield (key, lines)
73
74 if __name__ == '__main__':
75     start_time = datetime.now()
76     WordInverseIndex.run()
77     end_time = datetime.now()
78     elapsed_time = (end_time - start_time)*1000
79     sys.stderr.write("Total Seconds WordInverseIndex.py: ({0}) microseconds\n".format(elapsed_time.
80     total_seconds()))

```

---